

Predictive Model for Average Avocado Prices

December 16, 2020

Executive Summary:

Introduction:

Methods:

Exploratory Data Analysis:

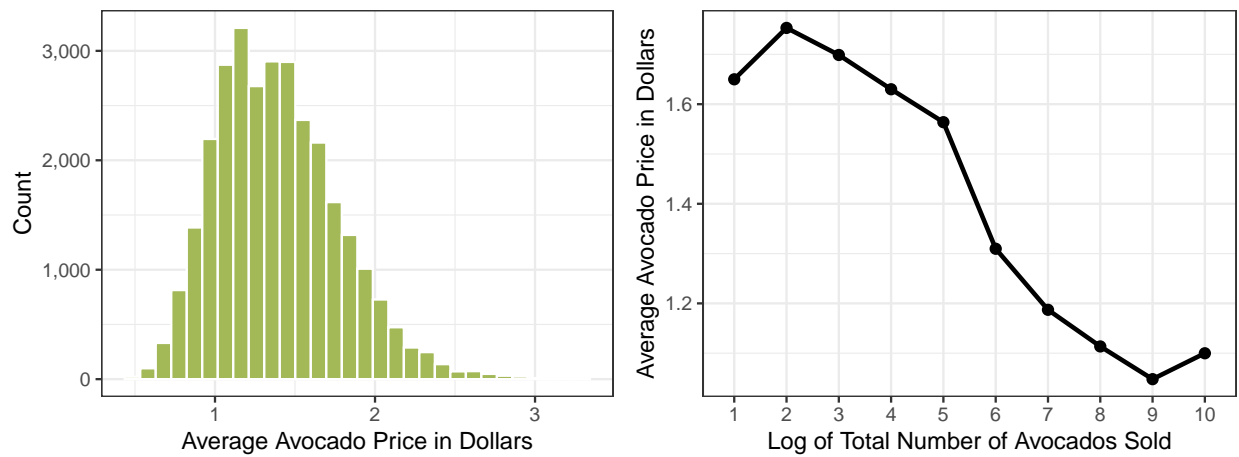


Figure 1

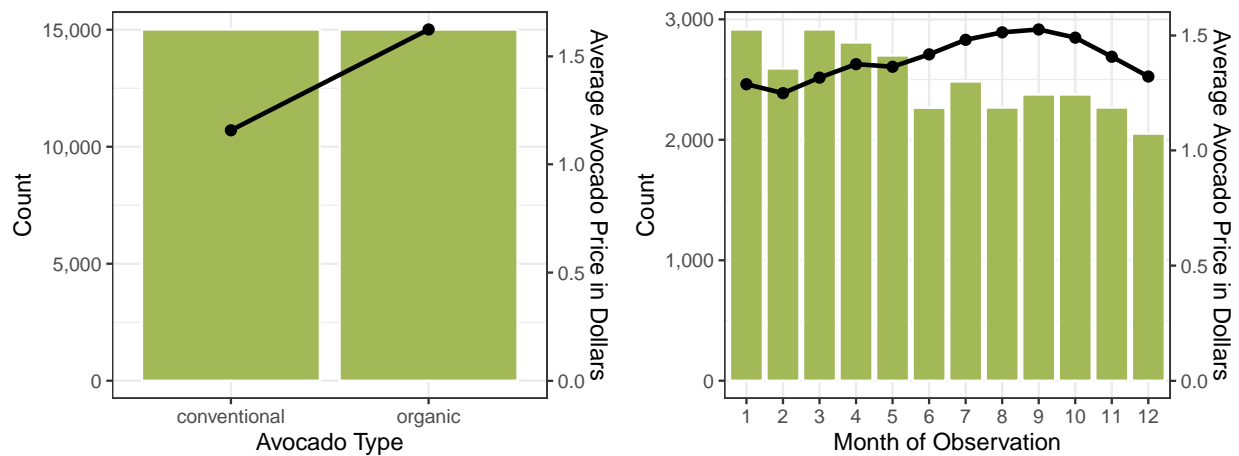


Figure 2




Figure 3




Figure 4

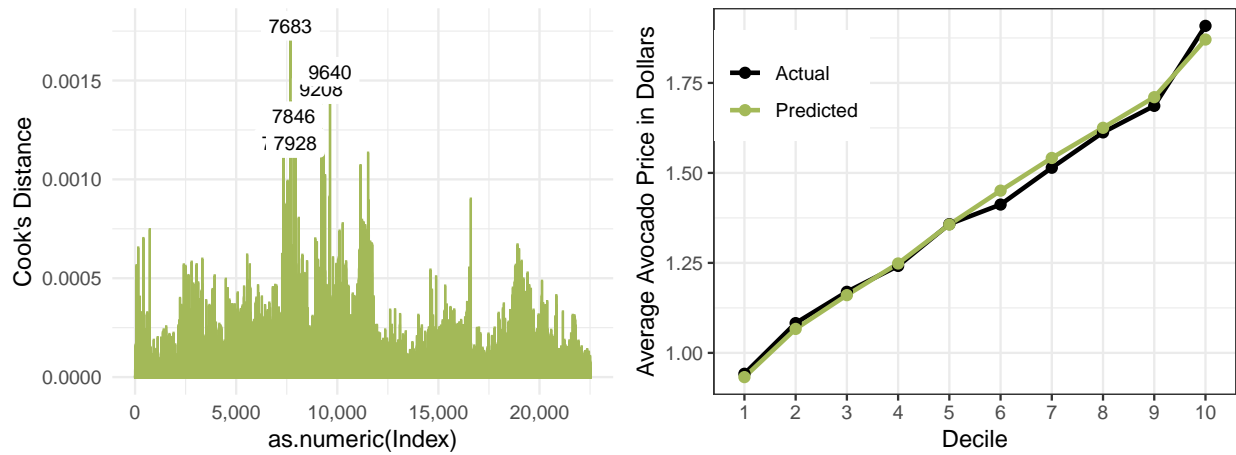


Figure 5

Term	Coef	SdError	F-Stat	pValue	2.5% CI	97.5% CI
(Intercept)	1.228	0.018	70.111	0.000	1.193	1.262
log_total_volume	-0.029	0.001	-23.237	0.000	-0.031	-0.026
month2	-0.037	0.008	-4.712	0.000	-0.052	-0.021
month3	0.031	0.008	4.066	0.000	0.016	0.045
month4	0.094	0.008	12.402	0.000	0.079	0.109
month5	0.080	0.008	10.502	0.000	0.065	0.095
month6	0.134	0.008	16.547	0.000	0.118	0.149
month7	0.198	0.008	25.137	0.000	0.182	0.213
month8	0.223	0.008	27.667	0.000	0.207	0.239
month9	0.242	0.008	30.527	0.000	0.227	0.258
month10	0.192	0.008	24.207	0.000	0.176	0.207
month11	0.113	0.008	13.961	0.000	0.097	0.129
month12	0.026	0.008	3.202	0.001	0.010	0.043
typeorganic	0.367	0.005	67.379	0.000	0.356	0.377
geography_bins2	0.160	0.005	33.397	0.000	0.150	0.169
geography_bins3	0.309	0.005	61.764	0.000	0.299	0.318
geography_bins4	0.563	0.008	69.124	0.000	0.547	0.579

Table 1: Summary regression of final model

Model Fitting/Inferences:**Conclusion:**

Appendix A: Supplemental Tables

Table 2: Summary Statistics for all numerical independent features

Statistic	N	Mean	St. Dev.	Min	Pctl(25)	Pctl(75)	Max
total_volume	30,021	939,255	3,813,519	85	14,299	489,803	63,716,144
4046	30,021	299,107	1,289,108	0	783	115,156	22,743,616
4225	30,021	284,901	1,169,078	0	2,814	140,947	20,470,573
4770	30,021	21,629	100,919	0	0	5,424	2,546,439
total_bags	30,021	333,534	1,415,618	0	8,374	159,174	31,689,189
small_bags	30,021	232,126	950,503	0	5,956	112,938	20,550,407
large_bags	30,021	95,185	467,210	0	352	36,068	13,327,601
xlarge_bags	30,021	6,223	38,137	0	0	560	1,022,564

	Model	Number of Features	MSE	Adj.R.squared	F.statistics	AIC
1	Initial Model	16.000	0.062	0.572	1879.873	1421.841
2	Stepwise Model	16.000	0.062	0.572	1879.873	1421.841
3	Model with Interaction Terms	78.000	0.060	0.588	413.437	598.507
4	Stepwise Model with Interaction Terms	77.000	0.060	0.588	418.808	597.053

Table 3: Regression validation metrics including MSE, R-squared adjusted, and AIC

	GVIF	Df	$GVIF^{1/(2 \cdot Df)}$
log_total_volume	2.713	1.000	1.647
month	1.007	11.000	1.000
type	2.673	1.000	1.635
geography_bins	1.031	3.000	1.005

Table 4

Appendix B: R Code

```

1 ##### Packages #####
2
3 library(magrittr)
4 library(tidymodels)
5 library(lubridate)
6 library(corrplot)
7 library(MASS)
8 library(broom)
9 library(car)
10 library(tidyverse)
11
12
13 ##### Parameters #####
14
15 avocado_color <- "#A3B958"
16
17
18 ##### Functions #####
19
20 ExploreVariable <- function(df, xvar, count = TRUE, x_axis){
21
22   group_df <- df %>%
23     group_by({{xvar}}) %>%
24     summarise(average_price = mean(average_price),
25               count = n(),
26               .groups = "drop")
27
28   ratio <- max(group_df$count) / max(group_df$average_price)
29
30   if(count == TRUE){
31     p <- ggplot(group_df, aes(x = factor({{xvar}}), group = 1)) +
32       geom_bar(aes(y = count), stat = "identity", fill = avocado_color, col = "white") +
33       geom_point(aes(y = average_price * ratio), size = 2, color = "black") +
34       geom_line(aes(y = average_price * ratio), size = 1, color = "black") +
35       scale_y_continuous(sec.axis = sec_axis(~./ratio, name = "Average Avocado Price
36         in Dollars"),
37                           label = scales::comma)
38
39     y_lab <- "Count"
40   } else {
41     p <- ggplot(group_df, aes(x = {{xvar}}, group = 1)) +
42       geom_point(aes(y = average_price), size = 2, color = "black") +
43       geom_line(aes(y = average_price), size = 1, color = "black")
44
45     y_lab <- "Average Avocado Price in Dollars"
46   }
47   p +
48     labs(y = y_lab, x = x_axis) +
49     theme_bw()
50 }
51
52
53 ValidationTable <- function(fit, model_type){

```

```

54 mod <- fit
55 fit_summary <- tibble(Model = model_type,
56                        "Number of Features" = length((coef(mod) %>% names())[1]),
57                        MSE = mean(mod$residuals^2),
58                        Adj.R.squared = summary(mod)$adj.r.squared,
59                        F.statistics = summary(mod)$fstatistic[[1]],
60                        AIC = AIC(mod))
61 return(fit_summary)
62 }
63
64 ##### Data #####
65
66 avocado <- read_csv("data/avocado-updated-2020.csv")
67
68
69 # descriptive statistics of the data
70
71 summary(avocado %>%
72         select_if(is.numeric))
73
74
75 # add deciles to continuous response
76
77 avocado %>%
78   mutate(
79     total_volume_bins = as_factor(cut(total_volume, breaks = 10,
80                                       include.lowest = TRUE, labels = FALSE)),
81     "4046_bins" = as_factor(cut(`4046`, breaks = 10,
82                                 include.lowest = TRUE, labels = FALSE)),
83     "4225_bins" = as_factor(cut(`4225`, breaks = 10,
84                                 include.lowest = TRUE, labels = FALSE)),
85     "4770_bins" = as_factor(cut(`4770`, breaks = 10,
86                                 include.lowest = TRUE, labels = FALSE)),
87     total_bags_bins = as_factor(cut(total_bags, breaks = 10,
88                                     include.lowest = TRUE, labels = FALSE)),
89     small_bags_bins = as_factor(cut(small_bags, breaks = 10,
90                                     include.lowest = TRUE, labels = FALSE)),
91     large_bags_bins = as_factor(cut(large_bags, breaks = 10,
92                                     include.lowest = TRUE, labels = FALSE)),
93     xlarge_bags_bins = as_factor(cut(xlarge_bags, breaks = 10,
94                                     include.lowest = TRUE, labels = FALSE)))
95
96 # add month
97
98 avocado %>%
99   mutate(month = factor(month(date)))
100
101
102 # feature engineer
103
104 price_by_location <- avocado %>%
105   group_by(geography) %>%
106   summarise(average_price = mean(average_price),
107             .groups = "drop") %>%
108   mutate(average_price_bins = cut(average_price, breaks = 4,
109                                   include.lowest = TRUE, labels = FALSE)) %>%

```

```

110   dplyr::select(geography, geography_bins = average_price_bins) %>%
111   distinct()
112
113 avocado %>%
114   left_join(price_by_location, by = "geography") %>%
115   modify_at("geography_bins", as_factor)
116
117
118 # log transformation of total_volume
119
120 avocado %>%
121   mutate(log_total_volume = log(total_volume)) %>%
122   mutate(log_total_volume_bins = as_factor(cut(log_total_volume, breaks = 10,
123                                               include.lowest = TRUE, labels = FALSE
124                                               )))
125
126 ##### Explore #####
127
128 # target variable: average_avocado price
129
130 ggplot(avocado, aes(average_price)) +
131   geom_histogram(fill = avocado_color, bins = 30, col = "white") +
132   theme_bw() +
133   labs(x = "Average Avocado Price in Dollars", y = "Count") +
134   scale_y_continuous(label = scales::comma)
135
136
137 # continuous predictors
138
139 ExploreVariable(avocado, total_volume_bins, count = FALSE,
140                x_axis = "Total Number of Avocados Sold")
141 ExploreVariable(avocado, log_total_volume_bins, count = FALSE,
142                x_axis = "Log of Total Number of Avocados Sold")
143
144 ExploreVariable(avocado, `4046_bins`, count = FALSE,
145                x_axis = "Total Number of Avocados with PLU 4046 Sold")
146 ExploreVariable(avocado, `4225_bins`, count = FALSE,
147                x_axis = "Total Number of Avocados with PLU 4225 Sold")
148 ExploreVariable(avocado, `4770_bins`, count = FALSE,
149                x_axis = "Total Number of Avocados with PLU 4770 Sold")
150 ExploreVariable(avocado, total_bags_bins, count = FALSE,
151                x_axis = "Total Number of Bags Sold")
152 ExploreVariable(avocado, small_bags_bins, count = FALSE,
153                x_axis = "Total Number of Small Bags Sold")
154 ExploreVariable(avocado, large_bags_bins, count = FALSE,
155                x_axis = "Total Number of Large Bags Sold")
156 ExploreVariable(avocado, xlarge_bags_bins, count = FALSE,
157                x_axis = "Total Number of Extra Large Bags Sold")
158
159 # categorical predictors
160
161 ExploreVariable(avocado, type, count = TRUE,
162                x_axis = "Avocado Type")
163 ExploreVariable(avocado, year, count = TRUE,
164                x_axis = "Year of Observation")

```

```

165 ExploreVariable(avocado, month, count = TRUE,
166                 x_axis = "Month of Observation")
167 ExploreVariable(avocado, geography, count = FALSE,
168                 x_axis = "Geography of The Avocado") +
169   coord_flip()
170 ExploreVariable(avocado, geography_bins, count = TRUE,
171                 x_axis = "Geography of the Avocado Group")
172
173
174 # correlation
175
176 corr_table <- avocado %>%
177   dplyr::select(total_volume, `4046`, `4225`, `4770`, total_bags, small_bags,
178                 large_bags, xlarge_bags, average_price) %>%
179   cor()
180
181 corr_table %>%
182   {.[order(abs(.[, 1]), decreasing = TRUE),
183     order(abs(.[, 1]), decreasing = TRUE)]} %>%
184   corplot(method = "number", type = "upper")
185
186
187 ##### Model Development #####
188
189 # select significant variables
190
191 avocado %>%
192   dplyr::select(average_price,
193                 log_total_volume,
194                 month,
195                 type,
196                 geography_bins)
197
198
199 # split data into train and test set
200
201 set.seed(123)
202
203 avocado_split <- initial_split(avocado, strata = average_price)
204
205 avocado_train <- training(avocado_split)
206
207 avocado_test <- testing(avocado_split)
208
209
210 # Initial model with all predictors
211
212 init_fit <- lm(average_price ~ .,
213               data = avocado_train)
214
215
216 # variable selection using stepAIC
217
218 step_fit <- stepAIC(init_fit, direction = "both", trace = FALSE)
219
220

```



```

221 # interaction
222
223 int_fit <- lm(average_price ~ .*.,
224             data = avocado_train)
225
226
227 # variable selection for interaction model
228
229 int_step_fit <- stepAIC(int_fit, direction = "both", trace = FALSE)
230
231
232 # generate iteration log
233
234 init_fit_summary <- ValidationTable(init_fit, "Initial Model")
235 step_fit_summary <- ValidationTable(step_fit, "Stepwise Model")
236 int_fit_summary <- ValidationTable(int_fit, "Model with Interaction Terms")
237 int_step_fit_summary <- ValidationTable(int_step_fit, "Stepwise Model with
    Interaction Terms")
238
239 bind_rows(init_fit_summary,
240           step_fit_summary,
241           int_fit_summary,
242           int_step_fit_summary) %>%
243   modify_if(is.numeric, round, 3)
244
245
246 # final model
247
248 final_fit <- init_fit
249
250 final_fit %>%
251   tidy() %>%
252   modify_if(is.numeric, round, 3)
253
254
255 ##### Model Diagnostic #####
256
257 # add rownames
258
259 avocado_train %>%
260   rownames_to_column()
261
262 # Residuals
263
264 avocado_train %>%
265   mutate(predict = predict(final_fit),
266          rstudent = rstudent(final_fit))
267
268
269 ## Influential Observations
270 ## Cook's D plot
271 ## identify D values > 4/(n-p-1) as a guide;
272 ## Cook and Weisberg recommend 0.5 and 1 (R uses these guides in default diagnostic
    plots below)
273
274 cutoff <- 4/((nrow(avocado_train) - length(final_fit$coefficients) - 2))

```

```

275
276 diag <- augment(final_fit) %>%
277   mutate(Index = 1:nrow(.))
278
279 diag %>%
280   mutate(high_cooksd = case_when(
281     .cooks_d > cutoff ~ 1, TRUE ~ 0),
282     col_stdresid = case_when(
283       .std.resid > 0 ~ 1,
284       .std.resid < 0 ~ 0),
285     high_hat = case_when(
286       .hat > .1 ~ 1,
287       TRUE ~ 0))
288
289 ## cook's distant ggplot
290
291 ggplot(diag, aes(x = as.numeric(Index), y = .cooks_d)) +
292   geom_bar(stat = "identity", col = avocado_color) +
293   labs(y = "Cook's Distance") +
294   theme_minimal() +
295   geom_label(data = diag %>% filter(.cooks_d > cutoff + .001),
296             aes(label = Index), label.size = NA, size = 3) +
297   scale_x_continuous(label = scales::comma)
298
299
300 ## Normality of Residuals
301
302 ggplot(diag, aes(x = .std.resid)) +
303   geom_histogram(bins = 30, col = "white", fill = avocado_color) +
304   theme_bw() +
305   labs(x = "Studentized Residuals",
306        y = "Count") +
307   scale_y_continuous(label = scales::comma)
308
309
310 # studentize residual plot
311
312 ggplot(diag, aes(x = Index, y = .std.resid)) +
313   geom_point(alpha = 0.2, col = avocado_color) +
314   labs(y = "Standard Residuals", x = "Index") +
315   theme_bw() +
316   theme(legend.position = "none") +
317   geom_label(data = diag %>% filter(.std.resid > 5 | .std.resid < -5),
318             aes(label = Index), label.size = NA, size = 3, hjust=-.25, vjust=.3) +
319   scale_x_continuous(label = scales::comma)
320
321
322 # outlier and influential points
323
324 diag %>%
325   filter(.std.resid > 5 | .std.resid < -5 | .cooks_d > cutoff + .001) %>%
326   dplyr::select(average_price, log_total_volume, month, type, geography_bins)
327
328
329 ## VIF
330 ## vif score seem very close to 1

```

```

331 vif(final_fit) # closer to 1 the better; 5-10 is moderate
332
333
334 ##### Predictions
335
336 # Account for sigma^2
337
338 sd_fit <- sd(final_fit$resid)
339
340
341 # predict
342
343 avocado_test %<>%
344   mutate(average_price_preds = predict(final_fit, newdata = .))
345
346
347 # errors
348
349 avocado_test %<>%
350   mutate(average_price_error = average_price - average_price_preds)
351
352 mse <- mean(avocado_test$average_price_error)^2
353
354
355 # create lift chart
356
357 avocado_test %<>%
358   mutate(average_price_decile = ntile(average_price_preds, n = 10))
359
360 decile_price <- avocado_test %<>%
361   group_by(average_price_decile) %>%
362   summarise(Actual = mean(average_price),
363             Predicted = mean(average_price_preds),
364             .groups = "drop") %>%
365   gather(key, price, -average_price_decile)
366
367 ggplot(decile_price,
368        aes(x = factor(average_price_decile), y = price,
369            group = key, color = key)) +
370   geom_line(size = 1) +
371   geom_point(size = 2) +
372   theme_bw() +
373   scale_color_manual(values = c("black", avocado_color)) +
374   labs(x = "Decile",
375        y = "Average Avocado Price in Dollars") +
376   theme(legend.title = element_blank(), legend.position = c(0.15,.8))

```

Listing 1: Appendix of Code